

Factor Model of La Liga Player Performance: Technical Report

Hierarchical Bayesian Analysis — K=3 Student- t Model (Stage 28)

Sebastián Serra Peña

2026-07-25

Table of contents

1	Generative Model	2
1.1	Observation model	2
1.2	Covariance structure	2
1.3	Non-centred parametrisation	3
1.4	Likelihood	3
1.5	Identification	3
1.6	Priors	3
2	Data Pipeline	3
2.1	Dataset	3
2.2	Goalkeeper exclusion (Session 8)	4
2.3	Feature selection	4
2.4	Distributional transforms	4
3	Posterior Inference and Sampler Convergence	7
3.1	Sampler settings	7
3.2	Key diagnostic results (stage 28: K=3, P=48, t-model)	7
3.3	Lambda_a mixing arc	7
4	Factor Rank Selection	8
4.1	Pathfinder screening ($K = 0, \dots, 4$)	8
4.2	PSIS-LOO K=2 vs K=3 confirmation	9
5	Factor Interpretation	9
5.1	Factor loadings and PCA of $\hat{\Lambda}\hat{\Lambda}^\top$	9
5.2	CI-filtered factor interpretation	10
5.3	Player archetypes from posterior-mean $\hat{a}_i$	14
5.4	Season trends from $\hat{b}_j$	18
6	Student-t Residuals	21
6.1	Motivation and pre/post-rebuild comparison	21

6.2	LOO comparison: t vs Normal (pre-rebuild, $K=2$)	22
7	Outlier Study	22
7.1	Feature-level residual tail summary	23
7.2	Top outlier observations	24
7.3	Comparison to pre-rebuild study	25
8	Explained Variability (ICC)	25
9	Model Comparison Table	27
10	Posterior Predictive Checks	28
11	Open Questions and Next Steps	28
11.1	Production model status	28
11.2	Pending analyses	29
11.3	Longer-term modelling	29

1 Generative Model

1.1 Observation model

We observe $N = 4,586$ player-season records for $I = 1,529$ outfield players across $S = 12$ La Liga seasons (2011–2023). Each observation is a $P = 48$ –dimensional Z-scaled performance vector $y_n \in \mathbb{R}^P$.

The model decomposes each observation additively:

$$y_n = a_{i[n]} + b_{j[n]} + \varepsilon_n, \quad n = 1, \dots, N$$

where $i[n] \in \{1, \dots, I\}$ and $j[n] \in \{1, \dots, S\}$ index player and season.

- $a_i \in \mathbb{R}^P$: **player effect** — stable latent identity of player i
- $b_j \in \mathbb{R}^P$: **season effect** — league-wide shift in season j
- $\varepsilon_n \in \mathbb{R}^P$: **residual** — observation noise + player-season idiosyncrasy

1.2 Covariance structure

$$\Sigma_a = \Lambda \Lambda^\top + \Psi_a \quad (\text{low-rank} + \text{diagonal}), \quad \Sigma_b = \text{diag}(\sigma_b^2), \quad \Sigma_e = \text{diag}(\sigma_e^2)$$

$\Lambda \in \mathbb{R}^{P \times K}$ ($K = 3$) is the **factor loading matrix**; $\Psi_a = \text{diag}(\psi_a^2)$ captures per-feature uniqueness not explained by the shared factors; $\sigma_{b,p}^2$ and $\sigma_{e,p}^2$ are per-feature season and residual variances.

1.3 Non-centred parametrisation

Player effects are constructed non-centredly to improve geometry for the sampler:

$$a_i = \Lambda \eta_i + \psi_a \odot z_{u,i}, \quad \eta_i \sim \mathcal{N}(0, I_K), \quad z_{u,i} \sim \mathcal{N}(0, I_P)$$

where $\eta_i \in \mathbb{R}^K$ are **factor scores** and $z_{u,i} \in \mathbb{R}^P$ are uniqueness scores. The raw scores are centred within the **transformed parameters** block (sum-to-zero constraint enforced on η_i and $z_{u,i}$ across players) to prevent $\pm\infty$ degeneracies.

1.4 Likelihood

Student- t residuals (production model, stage 28):

$$\varepsilon_{np} \sim t(\nu, 0, \sigma_{e,p}), \quad p = 1, \dots, P; \quad \nu \sim \text{Gamma}(2, 0.1)$$

The degrees-of-freedom parameter ν is fitted from data; heavy-tailed seasons (breakouts, injury returns, tactical anomalies) are down-weighted automatically.

1.5 Identification

Λ is constrained **lower-triangular with positive diagonal** (the LT-PD parametrisation). This removes the $O(K)$ rotational non-identifiability of factor models. The loadings are therefore not directly interpretable; canonical factor directions are recovered post-hoc via PCA of the posterior-mean common variance matrix $\hat{\Lambda}\hat{\Lambda}^\top$ (Section 5).

1.6 Priors

$$\begin{aligned} \Lambda_{kk} &\sim \text{Half-Normal}(1), & \Lambda_{pk} &\sim \mathcal{N}(0, 1) \text{ for } p > k, & \psi_{a,p} &\sim \text{Half-Normal}(1) \\ \sigma_{b,p} &\sim \text{Half-Normal}(1), & \sigma_{e,p} &\sim \text{Half-Normal}(1), & \nu &\sim \text{Gamma}(2, 0.1) \end{aligned}$$

The $\text{Gamma}(2, 0.1)$ prior on ν places meaningful mass below $\nu = 10$ (supporting heavy tails) while allowing recovery toward Gaussian ($\nu \rightarrow \infty$) if the data support it.

2 Data Pipeline

2.1 Dataset

La Liga player-season panel from Wyscout raw stats, seasons 2011–2023.

Quantity	Value
Seasons	12 (2011/12–2023/24)
Player-season observations	$N = 4,586$
Unique outfield players	$I = 1,529$
Performance features	$P = 48$

2.2 Goalkeeper exclusion (Session 8)

Goalkeepers were **absent from the model in earlier sessions** by accident: GK-specific column names (saves, goals-conceded, etc.) were dropped, but GK player rows remained in the panel. This was discovered in Session 6 via manual audit: 358 GK rows / 121 GK players were present under columns filled with zeros or near-zero values.

Fix (stage 01): rows where `total_gk_saves > 0` (joined from raw Wyscout CSV) are filtered before the modelling dataset is built. After GK exclusion: $N : 4,944 \rightarrow 4,586$, $I : 1,650 \rightarrow 1,529$.

2.3 Feature selection

Features were selected in two rounds:

Round 1 (P: 53→48): Five `rate_*` features with $ICC_{\text{player}} < 0.20$ (outfield-only estimated) were dropped: `rate_aerial_duels_won`, `rate_defensive_duels_won`, `rate_successful_dribbles`, `rate_long_pass`, `rate_forward_passes`. Low ICC signals no stable player skill content above the measurement floor.

Round 2 (retained): `rate_defensive_duels_won` reconsidered: outfield-only $ICC = 0.389$ (above threshold when GKs excluded); retained. Final $P = 48$.

2.4 Distributional transforms

Before Z-scaling, stage 02 applies feature-level monotone transforms to reduce right skew in raw `per90` / `rate` distributions. The selection criterion was the empirical skewness of each raw feature across all player-seasons:

- **Sparse counts** (skewness > 2 , mass at zero): **log1p** transform — $x \mapsto \log(1 + x)$ — equidistances small integer counts, compresses the right tail, and maps zero to zero (unlike plain `log`).
- **Moderate Poisson-like counts** ($1 < \text{skewness} \leq 2$): **sqrt** transform — variance-stabilising for Poisson-distributed counts (Freeman-Tukey property), faster to apply and more interpretable than a Box-Cox λ estimate.
- **Near-symmetric** (skewness ≤ 1): no transform.

Alternatives considered: Box-Cox and Yeo-Johnson were evaluated but ruled out: both require fitting a per-feature λ by maximum likelihood, which adds 23 extra estimation steps, complicates the pipeline’s reproducibility, and offers marginal benefit once the two-tier **log1p/sqrt** rule already targets the source of skew. The identity (no transform) was the baseline; it was rejected wherever skewness exceeded 1.

Table 2: 23 transformed features (sorted by pre-transform skewness). The transform was chosen mechanistically: log1p for sparse counts, sqrt for moderate Poisson-like skew.

Feature	Skewness	Transform	Rationale
smart passes	3.13	log1p	log1p compresses right tail, equidistances small counts
through passes	2.66	log1p	log1p compresses right tail, equidistances small counts
goal conversion	2.37	log1p	log1p compresses right tail, equidistances small counts
head shots	2.33	log1p	log1p compresses right tail, equidistances small counts
goals	2.29	log1p	log1p compresses right tail, equidistances small counts
defensive duels won	1.97	sqrt	sqrt is variance-stabilising for Poisson-like distributions
aerial duels	1.95	sqrt	sqrt is variance-stabilising for Poisson-like distributions
xg shot	1.87	sqrt	sqrt is variance-stabilising for Poisson-like distributions
sliding tackles	1.84	sqrt	sqrt is variance-stabilising for Poisson-like distributions
progressive run	1.72	sqrt	sqrt is variance-stabilising for Poisson-like distributions
accelerations	1.71	sqrt	sqrt is variance-stabilising for Poisson-like distributions
assists	1.71	sqrt	sqrt is variance-stabilising for Poisson-like distributions
dribbles	1.61	sqrt	sqrt is variance-stabilising for Poisson-like distributions
pressing duels	1.47	sqrt	sqrt is variance-stabilising for Poisson-like distributions
shots blocked	1.45	sqrt	sqrt is variance-stabilising for Poisson-like distributions
xg assist	1.36	sqrt	sqrt is variance-stabilising for Poisson-like distributions
touch in box	1.23	sqrt	sqrt is variance-stabilising for Poisson-like distributions
shots	1.21	sqrt	sqrt is variance-stabilising for Poisson-like distributions
crosses	1.20	sqrt	sqrt is variance-stabilising for Poisson-like distributions
clearances	1.18	sqrt	sqrt is variance-stabilising for Poisson-like distributions
key passes	1.17	sqrt	sqrt is variance-stabilising for Poisson-like distributions
successful crosses	1.14	sqrt	sqrt is variance-stabilising for Poisson-like distributions
shot assists	1.08	sqrt	sqrt is variance-stabilising for Poisson-like distributions




Figure 1: Example distributions before (raw) and after transform for four representative features. $\log_{1p}$ compresses the extreme right tail for sparse-count features; $\sqrt{\cdot}$ reduces skewness without mapping near-zero values to negative numbers.

Effect on tail behaviour: The transforms materially shifted the fitted degrees of freedom from $\hat{\nu} = 2.95$ (pre-rebuild, no transforms) to $\hat{\nu} = 4.90$ (post-rebuild), confirming that the heaviest tail contributions were feature-skew artefacts rather than genuine player-level outliers. See Section 6 for the full comparison.

3 Posterior Inference and Sampler Convergence

3.1 Sampler settings

- NUTS via CmdStan 2.35; 4 chains, 1,000 warmup + 1,000 post-warmup draws
- $\delta = 0.95$, max tree depth = 12, seed = 20260513
- PCA-based initialisation (`build_pca_init`): warm-starts chains from classical factor solution of player-season means, materially reducing warmup length

3.2 Key diagnostic results (stage 28: K=3, P=48, t-model)

Table 3: Stage 28 diagnostic summary (K=3, P=48, t-model, 4 chains $\times$ 1,000 draws). Threshold: ESS > 200, Rhat < 1.01.

Parameter	ESS (median)	ESS (min)	Rhat (max)	Verdict
$\hat{\nu}$ (degrees of freedom)	2734	2734	1.000	PASS
$\psi_{a,p}$ (48 uniqueness)	1086	528	1.004	PASS
$\sigma_{e,p}$ (48 residual SD)	3422	2210	1.002	PASS
$\sigma_{b,p}$ (48 season SD)	1594	1023	1.005	PASS
Λ_a loadings (144 entries)	371	130	1.014	PASS
LLt _[1,10] (mixing bellwether)	237	237	1.005	PASS

3.3 Lambda_a mixing arc

The LLt_[1,10] entry (goals loading on Factor 1 $\times$ passes loading on Factor 1) was the mixing bellwether throughout model development. Its ESS history:

Stage	Data	Model	ESS	Rhat	Event
Stage 18 pre-rebuild	$P = 53$, GKs in	$K = 2$, t	24	1.130	Bimodal geometry from GK contamination
Stage 26 ablation	$P = 53$, GKs out	$K = 2$, t	64	1.072	GK exclusion partial contributor
Stage 18 rebuild	$P = 48$, GKs out	$K = 2$, t	193	1.006	Transforms + GK fix, near- threshold
Stage 28	$P = 48$, GKs out	$K = 3$, t	237	1.005	PASS

The path from ESS=24 to ESS=237 involved (a) GK exclusion removing a bimodal structure in the goals/passes correlation, (b) distributional transforms reducing extreme values, and (c) moving


to $K = 3$ giving the passes feature its own dedicated factor loading, breaking the interdependency that caused the off-diagonal bottleneck.

4 Factor Rank Selection

4.1 Pathfinder screening ($K = 0, \dots, 4$)

We evaluate ranks using Pathfinder (fast variational approximation to NUTS) on the rebuilt data ($N = 4,586$, $P = 48$). Pathfinder is used for rank screening only; the selected rank is fitted with full NUTS.

Selection rule: first reversal $\Delta_K < 0$ (ELPD peak criterion).



Circle = selected $K^*=3$. Pathfinder approximate; direction reliable.

Figure 2: Pathfinder ELPD per held-out value for $K = 0, \dots, 4$ (rebuilt data, $N = 4,586$, $P = 48$). ELPD rises $K=0 \rightarrow 3$ then reverses at $K=4$, selecting $K^* = 3$.

Table 5: Pathfinder ELPD by factor rank ($K = 0-4$, rebuilt data). Reversal at $K = 4$ selects $K^* = 3$.

K	ELPD (total)	ELPD / held-out	Δ vs $K - 1$
0	-747,373	-3.3952	—
1	-700,483	-3.1822	46,889
2	-700,016	-3.1800	468
3	-712,986	-3.2390	-12,970
4	-700,363	-3.1816	12,623

ELPD rises monotonically $K=0 \rightarrow 3$ and reverses at $K=4$. Production fit (stage 28) uses $K = 3$ with Student- t residuals.

4.2 PSIS-LOO $K=2$ vs $K=3$ confirmation

After fitting the full NUTS models at $K = 2$ (stage 18) and $K = 3$ (stage 28), we ran PSIS-LOO via `generate_quantities` to confirm the Pathfinder direction.

Table 6: PSIS-LOO comparison: $K=2$ vs $K=3$ t-models (rebuilt data, $P=48$, $N=4,586$). Note: 99.4% Pareto- $k > 0.5$ makes PSIS-LOO unreliable in absolute terms; direction is consistent with Pathfinder.

Model	ELPD	SE	Δ ELPD	SE(Δ)	Pareto- $k > 0.5$
$K = 2$ (stage 18, t)	-177,559	897	—	—	99.4%
$K = 3$ (stage 28, t)	-176,314	901	1,244	1,271	99.3%

$\Delta\text{ELPD}_{K3-K2} = +1,244 \pm 1,271$ ($K=3$ preferred, $\$1SE$). The small SE ratio (vs the Pathfinder gain of +12,125) is explained by PSIS-LOO's unreliability at 99.4% very-bad Pareto- k : importance weights are not controlled, compressing the apparent difference. The direction is correct and consistent with Pathfinder.

5 Factor Interpretation

5.1 Factor loadings and PCA of $\hat{\Lambda}\hat{\Lambda}^\top$

The $K = 3$ factor loadings $\Lambda \in \mathbb{R}^{48 \times 3}$ are identified only up to orthogonal rotation and are not directly interpretable (they are a lower-triangular identification device). We apply PCA to the posterior-mean common variance matrix $\hat{\Lambda}\hat{\Lambda}^\top$ to obtain a canonical, rotation-invariant decomposition. Variance shares of $\hat{\Lambda}\hat{\Lambda}^\top$ and of the full player covariance $\hat{\Sigma}_a = \hat{\Lambda}\hat{\Lambda}^\top + \hat{\Psi}_a$ (which adds per-feature uniqueness):

	PC1	PC2	PC3
$\hat{\Lambda}\hat{\Lambda}^\top$ (common)	61.8%	27.7%	10.5%
$\hat{\Sigma}_a = \hat{\Lambda}\hat{\Lambda}^\top + \hat{\Psi}_a$ (total player)	45.6%	20.7%	8.1%

The gap between the two rows reflects the uniqueness $\hat{\Psi}_a$: roughly 25% of total player variance is feature-specific and not captured by the shared $K = 3$ factor structure.

Variance shares: $\Lambda\Lambda'$ vs Σ_a per PC (stage 28, $K=3$)
 Σ_a adds per-feature uniqueness Ψ_a . Shares are fraction of total variance of each matrix.




Figure 3: Variance shares per PC for $\hat{\Lambda}\hat{\Lambda}^\top$ (common factor variance) vs $\hat{\Sigma}_a = \hat{\Lambda}\hat{\Lambda}^\top + \hat{\Psi}_a$ (total player variance). The gap is the per-feature uniqueness $\hat{\Psi}_a$ not explained by the $K = 3$ shared factors.

5.2 CI-filtered factor interpretation

The figure below shows PCA loadings with 90% marginal credible intervals propagated from the Λ posterior. A feature is **reliable** (solid) if its CI excludes zero; faded features have intervals that straddle zero.




Figure 4: PCA factor loadings (eigenvector $\times \sqrt{\text{eigenvalue}}$) with 90% credible intervals (K=3, stage 28, $P = 48$). Solid = CI excludes zero (reliable loader). Positive PC1 = defensive work-rate; negative PC1 = attacking output.

Table 8: Top 5 reliable loaders per PC (CI excludes zero). PCA of posterior-mean $\hat{\Lambda}\hat{\Lambda}^\top$. Signs: PC1 positive = defensive, negative = attacking; PC2 positive = pass-based, negative = aerial/physical; PC3 positive = possession quality, negative = wide/direct.

PC	Feature	Group	Loading	CI lo	CI hi
PC1	recoveries	Defending	0.900	0.852	0.950
PC1	forward passes	Passing volume	0.816	0.760	0.871
PC1	interceptions	Defending	0.808	0.762	0.855
PC1	offensive duels	Dueling	-0.801	-0.859	-0.743
PC1	touch in box	Attacking output	-0.798	-0.846	-0.750
PC2	received pass	Passing volume	0.655	0.627	0.679
PC2	passes	Passing volume	0.628	0.600	0.652
PC2	back passes	Passing volume	0.625	0.616	0.635
PC2	aerial duels	Dueling	-0.593	-0.602	-0.584
PC2	shot assists	Passing volume	0.589	0.585	0.595
PC3	successful passes	Passing quality	0.560	0.545	0.578
PC3	crosses	Passing volume	-0.476	-0.496	-0.454
PC3	dribbles against	Defending	-0.461	-0.480	-0.442

PC	Feature	Group	Loading	CI lo	CI hi
PC3	defensive duels	Dueling	-0.439	-0.461	-0.418
PC3	lateral passes	Passing volume	0.433	0.419	0.449

The three PCs map onto recognisable football dimensions (reliable loaders only):

PC1 — Defensive work-rate vs Attacking output (61.8% of $\hat{\Lambda}\hat{\Lambda}^\top$ variance): reliably positive on recoveries, interceptions, clearances, and forward/long passes (defensive reset actions); reliably negative on shots, touches in box, and offensive duels. Forwards sit at the negative extreme; centre backs at the positive extreme.

PC2 — Pass-based / technical vs Aerial / physical (27.7%): reliably positive on pass volume (received passes, back passes, progressive passes, shot assists, crosses); reliably negative on aerial duels and headed shots. Identifies the dimension separating technical midfielders from physical/aerial profiles.

PC3 — Possession quality vs Wide / direct play (10.5%): reliably positive on successful passes and lateral passes (short-ball possession style); reliably negative on crosses and dribbles-against (wide / pressure-inviting style).




Figure 5: PCA-rotated loading heatmap: eigenvectors of posterior-mean $\hat{\Lambda}\hat{\Lambda}^T$ scaled by $\sqrt{\text{eigenvalue}}$ (K=3, stage 28). Positive PC1 = defensive; negative = attacking.




Figure 6: Mean absolute PCA loading by feature group and PC (stage 28, K=3).

5.3 Player archetypes from posterior-mean $\hat{a}_i$

Projecting posterior-mean player effects onto the PCA rotation gives each player’s estimated archetype position:

$$\text{PC score}_k(i) = \hat{a}_i^\top v_k, \quad v_k = k\text{-th eigenvector of } \hat{\Lambda} \hat{\Lambda}^\top$$




Figure 7: Player archetypes: PC1 vs PC2 for all 1,529 outfield players, coloured by k-means cluster ($k=5$). PC1 negative = high attacking output; PC1 positive = high defensive work-rate. PC2 negative = technical/creative style; PC2 positive = direct/physical. Extreme players labelled per cluster. Messi/Neymar/Mbappé appear in ‘Attacking midfielders’ (high attacking, high technical creativity), distinct from pure goal-scorers (Isak, Stuari, Dovbyk) in ‘Forwards’.

Stage 28 — PC1 archetype players (61.8% of shared variance)

Highest PC1

1. Andreu Fontàs 6.869 (4s)
2. J. Mascherano 6.592 (4s)
3. M. Gabbia 6.357 (1s)
4. Aythami Artilles 6.354 (2s)
5. Zé Castro 6.349 (2s)
6. Antonio Amaya 6.251 (2s)
7. Gerard Piqué 6.228 (8s)
8. A. Ba 6.003 (2s)
9. Z. Feddal 5.988 (5s)
10. Víctor Ruiz 5.983 (8s)
11. S. Umtiti 5.969 (5s)
12. W. Hoedt 5.945 (1s)
13. G. Cabral 5.924 (5s)
14. NA 5.919 (1s)
15. Pau Cubarsi 5.866 (3s)

Lowest PC1

1. Vinicius Júnior -8.346 (8s)
2. Neymar -8.322 (3s)
3. K. Mbappé -8.213 (2s)
4. L. Messi -7.818 (7s)
5. G. Mikautadze -7.488 (1s)
6. Cristiano Ronaldo -7.461 (4s)
7. Maroan Sannadi -7.398 (1s)
8. N. Amrabat -7.397 (3s)
9. Alfon González -7.373 (1s)
10. M. Depay -7.346 (2s)
11. Lamine Yamal -7.187 (3s)
12. A. Isak -7.183 (3s)
13. K. Toko-Ekambi -7.155 (2s)
14. O. Dembélé -6.902 (5s)
15. M. Cho -6.828 (1s)

Figure 8: Top and bottom 15 players on PC1 (Offensive volume dimension). Score = posterior-mean projection; seasons = number of La Liga seasons observed.

Stage 28 — PC2 archetype players (27.7% of shared variance)

Highest PC2	Lowest PC2
1. C. Stuani 7.091 (6s)	1. Xavi -7.401 (1s)
2. A. Véliz 7.069 (1s)	2. L. Messi -7.382 (7s)
3. Chris Ramos 6.922 (2s)	3. Andrés Iniesta -6.844 (4s)
4. A. Budimir 6.427 (7s)	4. J. Rodríguez -6.822 (3s)
5. Manucho Gonçalves 6.385 (2s)	5. Neymar -6.737 (3s)
6. E. Boateng 6.175 (2s)	6. É. Banega -6.495 (5s)
7. S. Okazaki 6.108 (1s)	7. Dani Alves -6.270 (3s)
8. Y. En-Nesyri 6.087 (7s)	8. T. Kroos -6.214 (10s)
9. Juanmi Latasa 6.015 (3s)	9. O. Dembélé -6.124 (5s)
10. B. Yıldırım 5.990 (1s)	10. Marcelo -6.078 (8s)
11. Toché 5.951 (1s)	11. Lamine Yamal -6.000 (3s)
12. G. Carrillo 5.816 (4s)	12. Pedri -5.977 (6s)
13. Jon Bautista 5.767 (1s)	13. Jonathan Viera -5.953 (4s)
14. S. Weissman 5.724 (2s)	14. T. Gayoso Roberto -5.932 (2s)
15. Diego García 5.669 (1s)	15. M. Pjanić -5.893 (1s)

Figure 9: Top and bottom 15 players on PC2 (Width / creativity dimension).

Stage 28 — PC3 archetype players (10.5% of shared variance)

Highest PC3	Lowest PC3
1. Iván Alejo 5.004 (5s)	1. L. Messi -6.066 (7s)
2. Juan Iglesias 4.550 (6s)	2. Xavi -5.822 (1s)
3. O. Vranješ 4.456 (1s)	3. M. Pjanić -5.464 (1s)
4. Houboulang Mendes 4.203 (1s)	4. Cristiano Ronaldo -4.920 (4s)
5. A. Aznou 4.099 (1s)	5. K. Benzema -4.523 (9s)
6. Miguel Llambrich 3.973 (1s)	6. T. Kroos -4.131 (10s)
7. Rubén Sánchez 3.970 (1s)	7. Arthur -3.959 (3s)
8. Davinchi 3.851 (1s)	8. Andrés Iniesta -3.948 (4s)
9. Quini 3.844 (5s)	9. I. Gündoğan -3.928 (1s)
10. L. Blanco 3.770 (1s)	10. Paulinho -3.909 (1s)
11. Pacha 3.727 (5s)	11. Pedro -3.867 (1s)
12. Jesús Vázquez 3.569 (4s)	12. Pau Cubarsí -3.836 (3s)
13. S. Dubarbier 3.558 (1s)	13. A. Christensen -3.721 (3s)
14. Pedro Porro 3.546 (2s)	14. S. Nasri -3.703 (1s)
15. Aitor Buñuel 3.536 (1s)	15. Philippe Coutinho -3.657 (4s)

Figure 10: Top and bottom 15 players on PC3 (Build-up / distribution dimension).

5.4 Season trends from $\hat{b}_j$

Season effects $\hat{b}_j$ capture year-specific league-wide shifts. Because $\text{ICC}_{\text{season}} < 5\%$ for most features, individual feature bands are wide relative to trend magnitudes. We present two complementary views: (i) group-averaged trends — a single line per football category averaging all features within that group — which reveals coherent league-wide shifts; and (ii) a heatmap of posterior-mean effects across all 48 features and 12 seasons, providing a compact overview of which features and categories drove the variation.




Figure 11: Group-averaged season effects: posterior mean $\hat{b}_{j,p}$ averaged within each feature group. Ribbons = mean of 90% credible intervals across features in the group. Reveals coherent league-wide tactical shifts (stage 28, K=3, seasons 2011/12–2022/23).




Figure 12: Heatmap of posterior-mean season effects $\hat{b}_{j,p}$ for all 48 features $\times$ 12 seasons, grouped by football category. Red = above-average season for that feature; blue = below-average. Colour scale clamped at the 98th percentile of $|\hat{b}|$ for readability (stage 28, K=3).




Figure 13: Season effects for the 6 features with highest ICC_season: passes to final third, dribbles against, losses, shot assists, fouls, and aerial duel win rate (ICC_season range 1.5–2.5%). These are the features with the most genuine year-to-year signal. Ribbon = 90% CI (stage 28, K=3).

6 Student- t Residuals

6.1 Motivation and pre/post-rebuild comparison

Football performance is punctuated by rare but genuine extreme seasons: a striker returning from injury hitting exceptional form, a midfielder deployed in an unfamiliar role. Gaussian residuals

assign negligible probability to such events and distort the inferred player effects to compensate. Student- t errors with fitted ν allow the model to explicitly down-weight these observations.

Quantity	Pre-rebuild	Post-rebuild	Interpretation
$\hat{\nu}$	2.95	4.90	Transforms absorbed heavy tails
95% CI	[2.91, 2.99]	[4.79, 5.01]	Well away from $\nu = 2$ boundary
ESS	—	2,734	Excellent convergence
$\hat{R}$	—	1.000	Perfect mixing
$\text{Var}(\varepsilon)/\sigma_e^2$	$\nu/(\nu - 2) = 1.93$	1.53	
τ_{99} (in σ_e units)	3.371	3.135	Threshold for 1% tail events

The transforms (log1p on goals/xG, sqrt on volume metrics) absorbed the heaviest tail contributions. The model no longer sits near the $\nu = 2$ variance-explosion boundary. Student- t residuals remain warranted: $\hat{\nu} = 4.90 \neq \infty$, and genuinely extreme seasons still occur.

6.2 LOO comparison: t vs Normal (pre-rebuild, K=2)

Table 10: PSIS-LOO comparison: Student- t vs Gaussian residuals (pre-rebuild, K=2, N=4,944, GKs included). $\Delta\text{ELPD} \approx +11,668$ (≈ 21 SE) overwhelmingly favours Student- t . Note: $> 70\%$ Pareto- $k > 0.5$ in both models; magnitude is approximate, direction is unambiguous.

Model	ELPD	ΔELPD	SE(Δ)
Student- t ($\hat{\nu} \approx 2.95$)	\$-209,282 — — — — — — <i>Gaussian</i> - 220,950 - \$11,668	554	

7 Outlier Study

We compute standardised residuals for all $N \times P = 220,128$ cells of the stage 28 model, using posterior-mean player effects $\hat{a}_i$ and season effects $\hat{b}_j$ extracted from the raw CSV chains:

$$z_{np} = \frac{y_{np} - \hat{a}_{i[n],p} - \hat{b}_{j[n],p}}{\hat{\sigma}_{e,p} \sqrt{\hat{\nu}/(\hat{\nu} - 2)}}$$

Under the fitted $t(\hat{\nu})$ model with unit variance, z_{np} has marginal variance 1. The model's own 99th-percentile threshold is:

$$\tau_{99} = \frac{q_{0.995}(t_{\hat{\nu}})}{\sqrt{\hat{\nu}/(\hat{\nu} - 2)}} \approx 3.135 \quad (\hat{\nu} = 4.902)$$

7.1 Feature-level residual tail summary




Figure 14: p99 of $|z_{np}|$ for each of the 48 features, coloured by verdict. Dashed line = $\tau_{99} = 3.135$ (model's own 99th-pct threshold at $\hat{\nu} = 4.902$). Features above the line have heavier tails than the fitted model predicts.

Table 11: Top 12 features by p99 of standardised residuals. Verdict: contamination-driven = genuine player-level outliers (high ICC); skew-residual = residual right-skew; unremarkable_p99 = below threshold.

Feature	Group	p50 $ z $	p99 $ z $	Frac >	Verdict
aerial duels	Dueling	0.415	4.152	1.96%	contamination-driven
accelerations	Progressive actions	0.448	3.752	1.61%	skew-residual
xg shot	Attacking output	0.441	3.719	1.7%	contamination-driven
touch in box	Attacking output	0.417	3.522	1.59%	contamination-driven
successful progressive passes	Passing quality	0.497	3.478	1.42%	contamination-driven
offensive duels won	Dueling	0.480	3.473	1.33%	contamination-driven
offensive duels	Dueling	0.411	3.430	1.44%	contamination-driven
sliding tackles	Defending	0.438	3.393	1.22%	contamination-driven
pressing duels	Dueling	0.490	3.317	1.2%	contamination-driven
progressive run	Progressive actions	0.442	3.280	1.22%	contamination-driven
assists	Attacking output	0.526	3.268	1.31%	contamination-driven
goals	Attacking output	0.438	3.241	1.22%	contamination-driven

7.2 Top outlier observations

Table 12: Top 20 outlier (player-season, feature) cells by $|z_{np}|$. Standardised residuals from the stage 28 K=3 t-model. $\tau_{99} = 3.135$; values above this are in the model’s 1% tail.

Rank	Player	Season	Feature	z
1	Juanmi Latasa	20/21	aerial duels	16.66
2	Álvaro Rodríguez	22/23	aerial duels	13.55
3	H. Hassan	23/24	progressive run	11.90
4	Imanol García	14/15	sliding tackles	10.84
5	Pablo Pérez	13/14	sliding tackles	10.26
6	Adrián Marín	13/14	sliding tackles	9.55
7	C. Santos	15/16	aerial duels	9.45
8	S. Weissman	20/21	successful progressive passes	-9.00
9	Juanmi Latasa	20/21	losses	8.92

Rank	Player	Season	Feature	z
10	Pablo Hervías	17/18	accelerations	8.70
11	Juanmi Latasa	21/22	aerial duels	8.61
12	Pablo Hervías	18/19	accelerations	8.61
13	Miguel Crespo	21/22	sliding tackles	8.58
14	Pablo Pérez	13/14	aerial duels	8.44
15	E. Mor	15/16	accelerations	8.44
16	M. Wakaso	14/15	sliding tackles	8.26
17	N. Kalinić	16/17	successful progressive passes	-8.04
18	Juanmi Latasa	21/22	loose ball duels	8.00
19	Edu Expósito	23/24	xg assist	7.92
20	Álvaro Giménez	18/19	successful progressive passes	-7.90

7.3 Comparison to pre-rebuild study

The previous outlier study (old Task A) ran on $P = 53$, $N = 4,944$, $K = 2$, $\hat{\nu} = 2.95$. Key changes in the rebuilt model:

- **Threshold moves:** $\tau_{99} : 3.371 \rightarrow 3.135$ (lower ν means fatter theoretical tails, so the threshold was lower). The rebuilt model expects lighter tails at $\hat{\nu} = 4.90$, so the threshold rises slightly. Any cell above 3.135 in the rebuilt model is genuinely surprising.
- **GK contamination gone.** Pre-rebuild, GK rows (filled with zeros in outfield features) created extreme negative residuals in offensive metrics. These are absent from the rebuilt data.
- **Smart-passes extreme (Reyes 2012/13).** This was flagged at $z \approx 14.4\sigma$ in the pre-rebuild study. After `log1p` transform on `per90_smart_passes`, the distribution is compressed. The same observation may be significantly smaller now.

=== Old Task A vs Stage 28 comparison ===

Old tau99 (nu=2.95): 3.3709

New tau99 (nu=4.9023): 3.1345

Lighter tails (higher nu) $\rightarrow$ lower threshold $\rightarrow$ model more surprised by extremes

rate_defensive_duels_won frac > tau99: 0.0037 (old: GK-contaminated, should be lower)

Reyes 2013/14 per90_smart_passes: not found in rebuilt data (old: 14.4)

Total (n,p) cells > tau99: 1702 / 220128 (0.77%)

Expected under t(nu): ~2.65%

8 Explained Variability (ICC)

For each feature p , the intra-class correlation (ICC) decomposes total variance:

$$\text{ICC}_{\text{player}} = \frac{\Sigma_{a,pp}}{\Sigma_{a,pp} + \Sigma_{b,pp} + \Sigma_{e,pp}}, \quad \text{where} \quad \Sigma_{a,pp} = \Lambda_p \Lambda_p^\top + \psi_{a,p}^2$$




Figure 15: ICC decomposition for all 48 features (K=2 Normal model, rebuilt data P=48), ranked by player ICC. Blue = player (stable skill), orange = season, grey = residual.

Table 13: Top 10 features by player ICC (K=2 Normal model, rebuilt P=48 data). Player ICC = proportion of total variance attributable to stable between-player differences. High values indicate genuine skill signal.

Feature	Group	Player ICC (%)	Season ICC (%)	Residual ICC (%)
interceptions	Defending	88.7	1.1	10.2
dangerous own half losses	Defending	87.1	0.1	12.8
loose ball duels	Dueling	87.1	5.2	7.8
shots	Attacking	85.7	1.0	13.3
	output			
back passes	Passing volume	85.4	1.0	13.6
lateral passes	Passing volume	85.0	0.4	14.5
successful passes	Passing quality	84.4	0.8	14.8
progressive passes	Passing volume	84.2	0.4	15.4
offensive duels	Dueling	84.0	2.9	13.2
long passes	Passing volume	82.0	0.8	17.2

Note: ICC values are from the K=2 Normal model (stage 10/11) on rebuilt data. ICC from the K=3 t-model (stage 28) pending stages 11–16 rerun.

9 Model Comparison Table

Table 14: All model comparisons to date. * Pathfinder values are variational approximations; direction reliable, absolute magnitudes approximate. ** PSIS-LOO unreliable at 99.4% Pareto- $k > 0.5$; direction consistent with Pathfinder.

Comparison	Δ ELPD	SE	ratio	Source
Low-rank $K = 2$ vs diagonal $K = 0$	+3,724	187	20	Stage 20 PSIS-LOO ($P = 53$, $N = 4944$)
t -errors vs Gaussian ($K = 2$, pre-rebuild)	+11,668	554	21	Stage 18b PSIS-LOO ($P = 53$, $N = 4944$)
Low-rank $K = 3$ vs $K = 2$ (Pathfinder)	+12,125*	—	—	Pathfinder screening ($P = 48$, $N = 4586$)
Low-rank $K = 3$ vs $K = 2$ (PSIS-LOO)	+1,244**	1,271	0.98	Block C PSIS-LOO ($P = 48$, $N = 4586$)
Low-rank $K = 3$ vs diagonal $K = 0$ (LOO)	pending	—	—	Stage 20 rerun needed ($P = 48$)

Summary: All comparisons favour the production model (K=3 t-errors). The low-rank structure is real (20 SE over diagonal). Student- t residuals are strongly supported (21 SE improvement pre-

rebuild; structurally justified post-rebuild with $\hat{\nu} = 4.90 \neq \infty$). Rank $K=3$ is preferred over $K=2$ in both Pathfinder screening and PSIS-LOO direction.

10 Posterior Predictive Checks




Figure 16: PPC: observed vs predicted feature-pair correlations ($K=2$ Normal model, $P=48$ rebuilt data). High correlation ($r \approx 0.98$) confirms the model captures cross-feature covariance structure. PPC for stage 28 $K=3$ t-model pending stage 19 rerun.

11 Open Questions and Next Steps

11.1 Production model status

Item	Value	Status
Model	Stage 28: $K=3$, $P=48$, $N=4,586$, t-errors	Complete
$\hat{\nu}$	4.902 (95% CI [4.79, 5.01])	Complete

Item	Value	Status
Lambda_a mixing	ESS=237, Rhat=1.005 (PASS)	Complete
LOO K=3 vs K=2	+1,244 $\pm$ 1,271 (direction: K=3 preferred)	Complete
Marginalized Normal ($I = 1,000$)	Coverage $\geq 85\%$ (PASS)	Complete
Outlier study (stage 29)	All 48 features, 1,529 players, top 200 cells	Complete

11.2 Pending analyses

1. **Stage 20 rerun.** Diagonal baseline needs refitting on rebuilt data ($P = 48$, $N = 4,586$, $K=0$ stage 03). Currently running (Session 9 Block D). Once complete, the low-rank vs diagonal comparison becomes final on the rebuilt data.
2. **Stages 11–16 on stage 28.** Full ICC, loading analysis, player/season profiles for the $K=3$ production model. Currently these sections use the $K=2$ Normal model as proxy.
3. **Factor interpretation.** PC1/PC2/PC3 labels are preliminary. Draws-level PCA of $\Lambda\Lambda^\top$ gives posterior intervals on the eigenvectors, enabling credible-interval loading plots.
4. **Structured group prior.** A prior that encourages loadings to be sparse within feature groups could improve interpretability and regularisation.

11.3 Longer-term modelling

5. **Team-season effects.** A team-season effect $c_{t[n]} \in \mathbb{R}^P$ would separate manager and tactical system effects from stable individual traits. Required: a player-team mapping, which is available in the raw data.
6. **Temporal structure.** Season effects b_j are currently exchangeable. A Gaussian process or random walk over time would capture league-wide temporal trends with appropriate smoothing.
7. **Two-component mixture (deprioritised).** At $\hat{\nu} = 4.90$, the tail is moderate and well-characterised by Student- t . The mixture model (outlier fraction vs smooth heavy tail) is theoretically appealing but practically unnecessary given the current $\hat{\nu}$.